



Interpretation of the plot

- **X-axis:** Actual delay, ground truth
- **Y-axis:** Predicted delay from a linear model
- The dashed line is the identity line (perfect prediction)

Positives

- As actual delay increases, predictions increase slightly
- Prediction variance is reasonably tight

Negatives

- For `actual_delay > 6` predictions cluster near 5–6 min, not growing linearly
- For `actual_delay < 2` predictions cluster near 3-4 min, too high

This means we are biased towards the mean

Model coefficients

0	hour	0.065893
1	dayofweek	-0.010351

```
2 is_weekend      -0.263213
3 direction_Linden 1.416742
```

- **direction_Linden = 1.41**: Linden-bound trains are consistently ~1.4 min more delayed than Howard-bound.
- **is_weekend = -0.26**: Weekend trains are more punctual. This might be due to the lack of data.
- **hour** and **dayofweek**: Effect is minimal, suggests that these have no effect.